

DATA

SATURDAYS

MADRID



Delta Lake Tables 101

30/11/2024

15:25-16:15

Kamil Nowinski



Sponsors:



#DataSatMadrid

¡Gracias sponsors!

Platinum



Gold



Silver



Thanks to all the Sponsors

Final Quiz + Raffle + Data Beers included (18:30h)

Stay for the Fun Quiz
and enter the Prize Draw
(only for attendees)
And much more...!



Sponsor
session/stand only



Sponsor
session/stand only



Sponsor stand only



Amby.net



Kamil Nowiński



Microsoft Data Platform **MVP**, **Databricks Champion**
Speaker, blogger, **YouTuber**, data enthusiast

Group Manager at [Avanade](#) UK&I
>20 years of experience in IT

Founder of blog [AzurePlayer](#)

GitHub: **#adftools**, SCD Merge Wizard and more...

Member of the Data Community PL

Former co-organiser of SQLDay (PL), Data Relay (UK),
volunteer at SQLBits (UK)

SQL Server Certificates:

MCITP, MCP, MCTS, MCSA, MCSE Data Platform,
MCSE Data Management & Analytics, DevOps Expert

Moreover: Bicycle, Running, Digital photography
[@AzurePlayer.bsky.social](#)

Blog

- Technical posts
- Various skill level
- Cheat sheets
- Recommended books
- Many useful other links
- Interviews (Podcast)
- YouTube Channel:
www.AzurePlayer.net/YouTube



Azure Player

Play with data & have fun!

www.AzurePlayer.net





Ask SQL Family ▶ Play all

This set of videos promote episodes of "Ask SQL Family" podcast.



Query data in a KQL database in Microsoft Fabric | Lab 12

Kamil Data Geek - Azure explained • 368 views • 2 months ago



Get started with Data Activator in Microsoft Fabric | Lab 11

Kamil Data Geek - Azure explained • 705 views • 2 months ago



Ingest data with Spark and Microsoft Fabric notebooks | Lab 10

Kamil Data Geek - Azure explained • 498 views • 3 months ago



Use real time eventstreams in Microsoft Fabric | Lab 09

Kamil Data Geek - Azure explained • 335 views • 3 months ago



Pawel Potasinski - What's important in your public...

Kamil Data Geek - Azure explained
54 views • 8 days ago



Simon Whiteley - What is hiding behind the term...

Kamil Data Geek - Azure explained
98 views • 2 months ago



Mladen Prajdić - Convince your boss to upgrade SQL...

Kamil Data Geek - Azure explained
37 views • 2 months ago



Alex Whittles - Successful BI business without autopilot |...

Kamil Data Geek - Azure explained
35 views • 3 months ago



Chris Webb - future of SSAS On-Premise and a few more...

Kamil Data Geek - Azure explained
140 views • 3 months ago



Kalen Delaney - story about @sqlqueen twitter handle |...

Kamil Data Geek - Azure explained
24 views • 4 months ago

www.AzurePlayer.net/YouTube

Slides available:

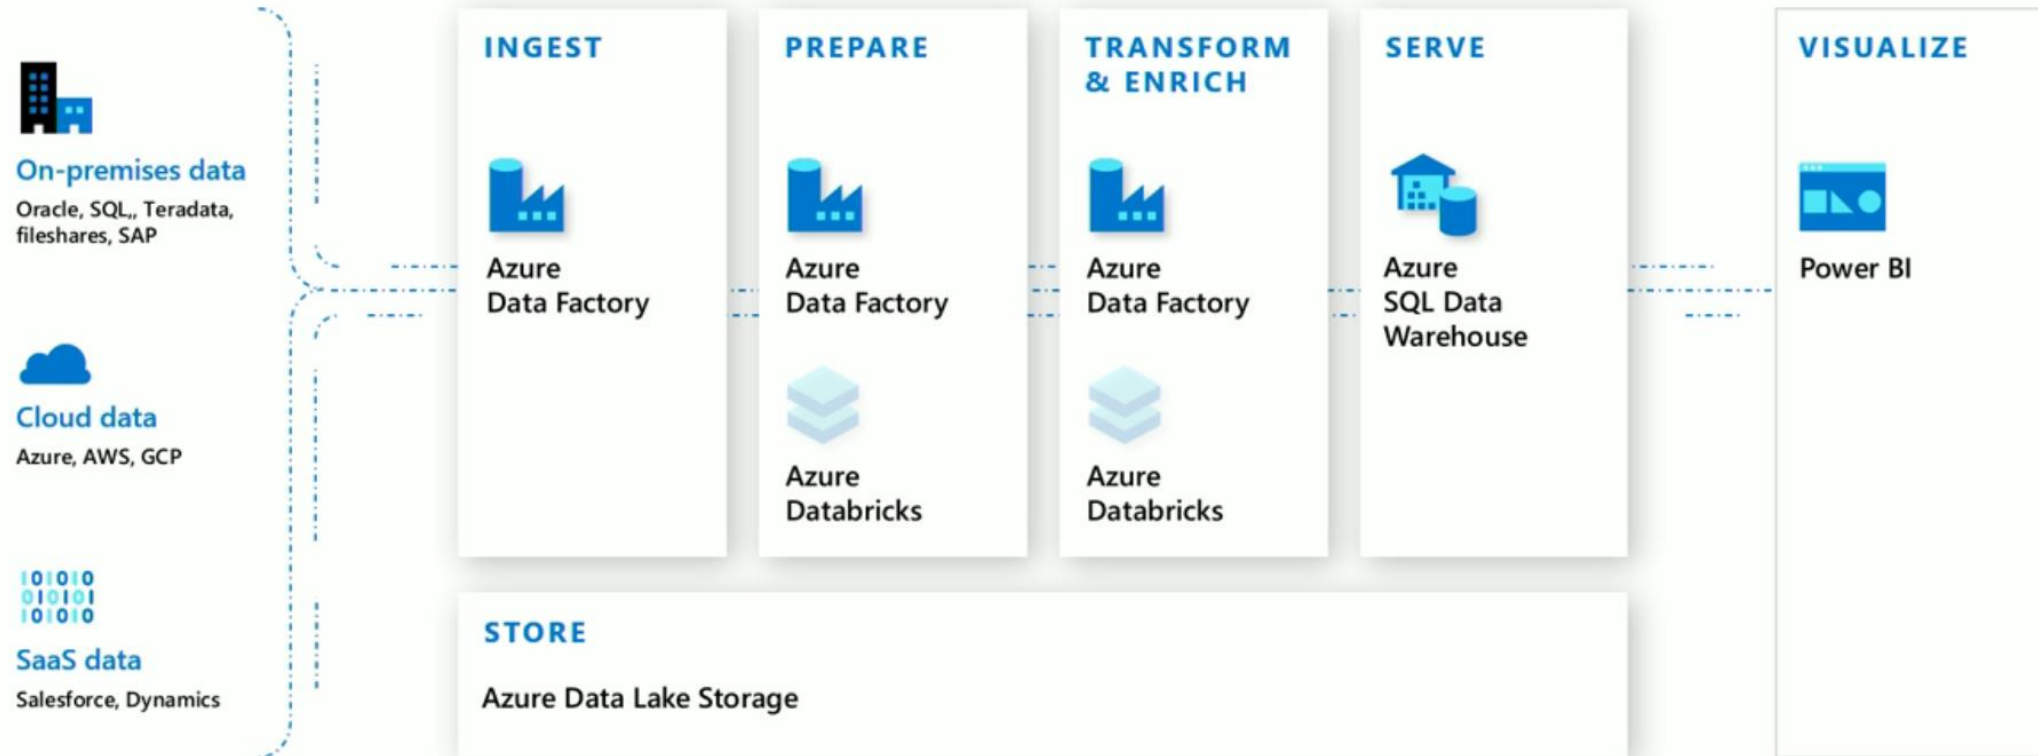


<https://azureplayer.net/slides>


AGENDA

- A bit of theory...
- ... and history
- Key features of Delta Lake Tables
- Let's practice: DEMO time
 - Delta Lake Tables in Lake Databases (Databricks)
- Key takeaways

Modern Data Warehouse




Data Lake + Warehouse = Lakehouse



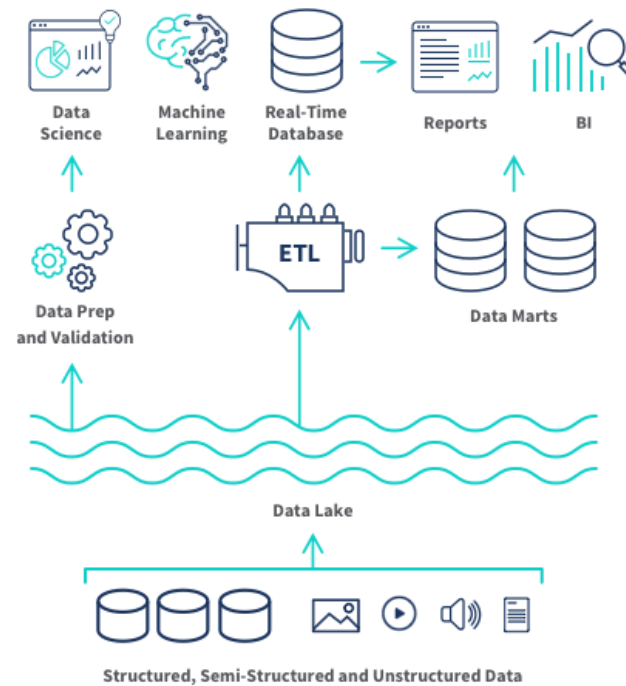
> [What Is a Lakehouse?](#)

Data Warehouse evolution

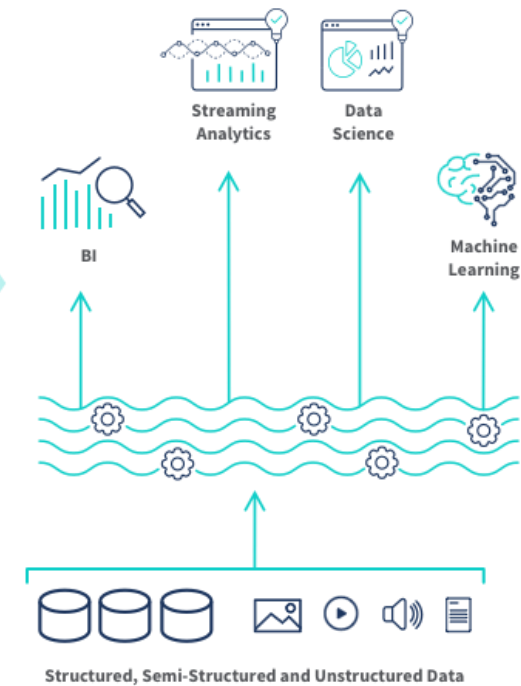
LATE 1980'S
Data Warehouse



2011
Data Lake



2020
Lakehouse



Parquet format

What is Parquet?

Apache Parquet is an **open source, column-oriented** data file format designed for efficient data storage and retrieval. It provides efficient data compression and encoding schemes with enhanced performance to handle complex data in bulk. Apache Parquet is designed to be a common interchange format for both batch and interactive workloads.

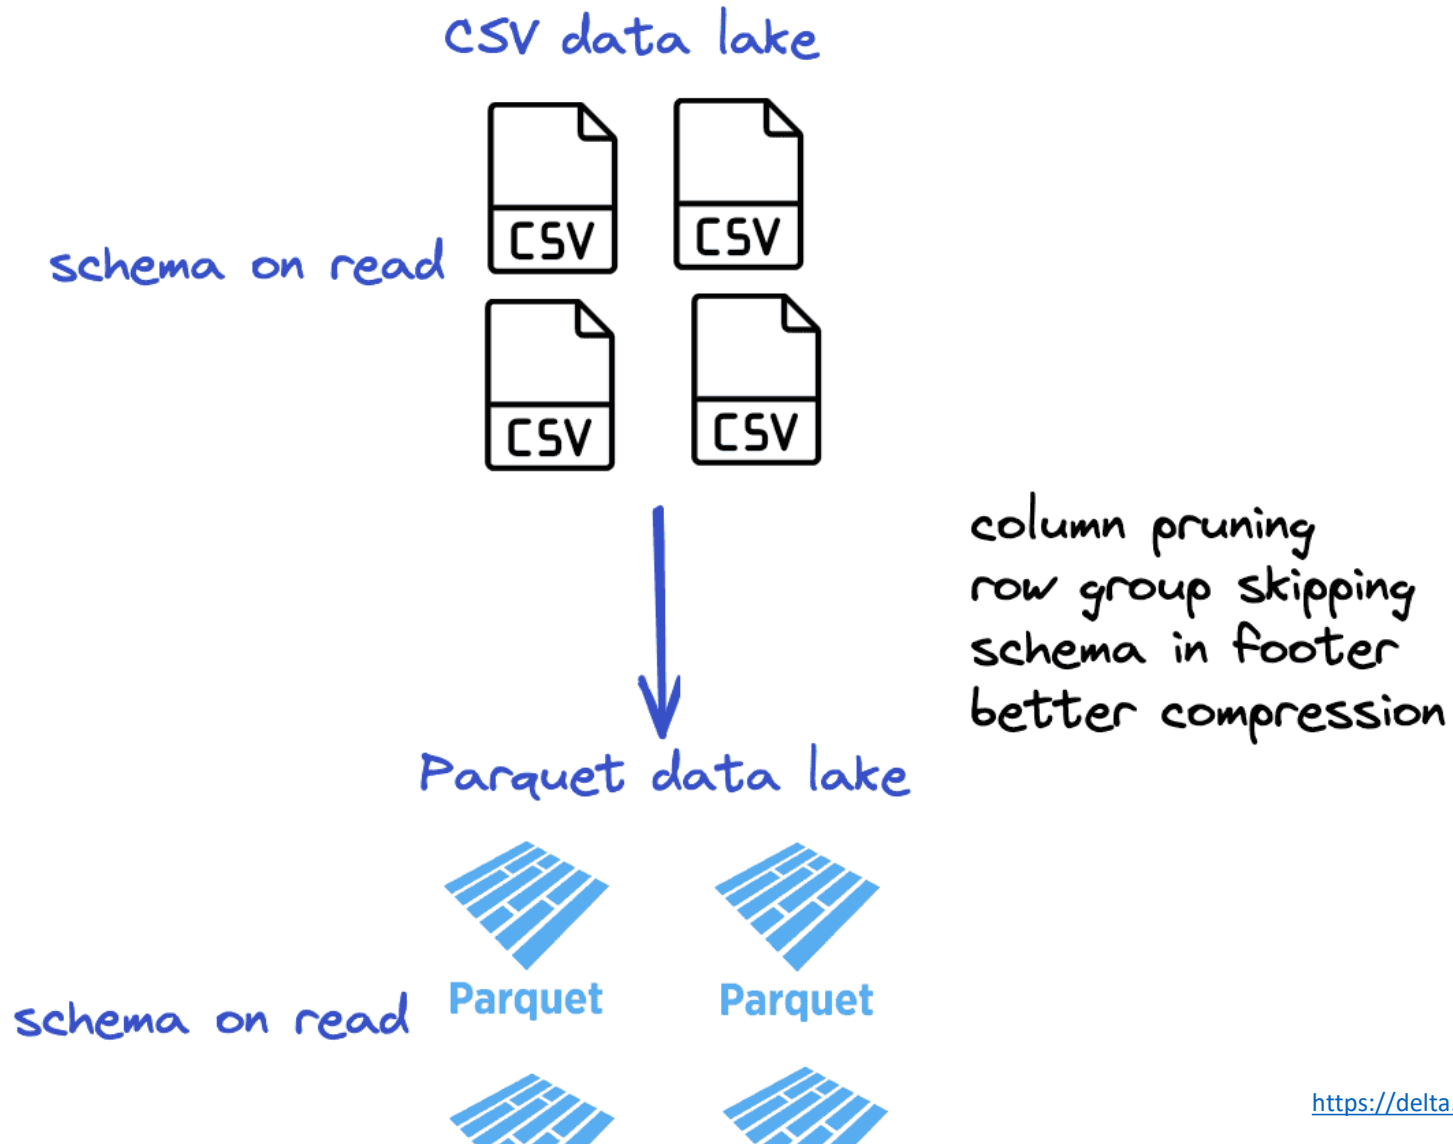
Characteristics of Parquet:

- Free and open-source file format
- Language agnostic
- Column-based format
- Used for analytics (OLAP) use cases
- Highly efficient data compression and decompression
- Supports complex data types and advanced nested data structures

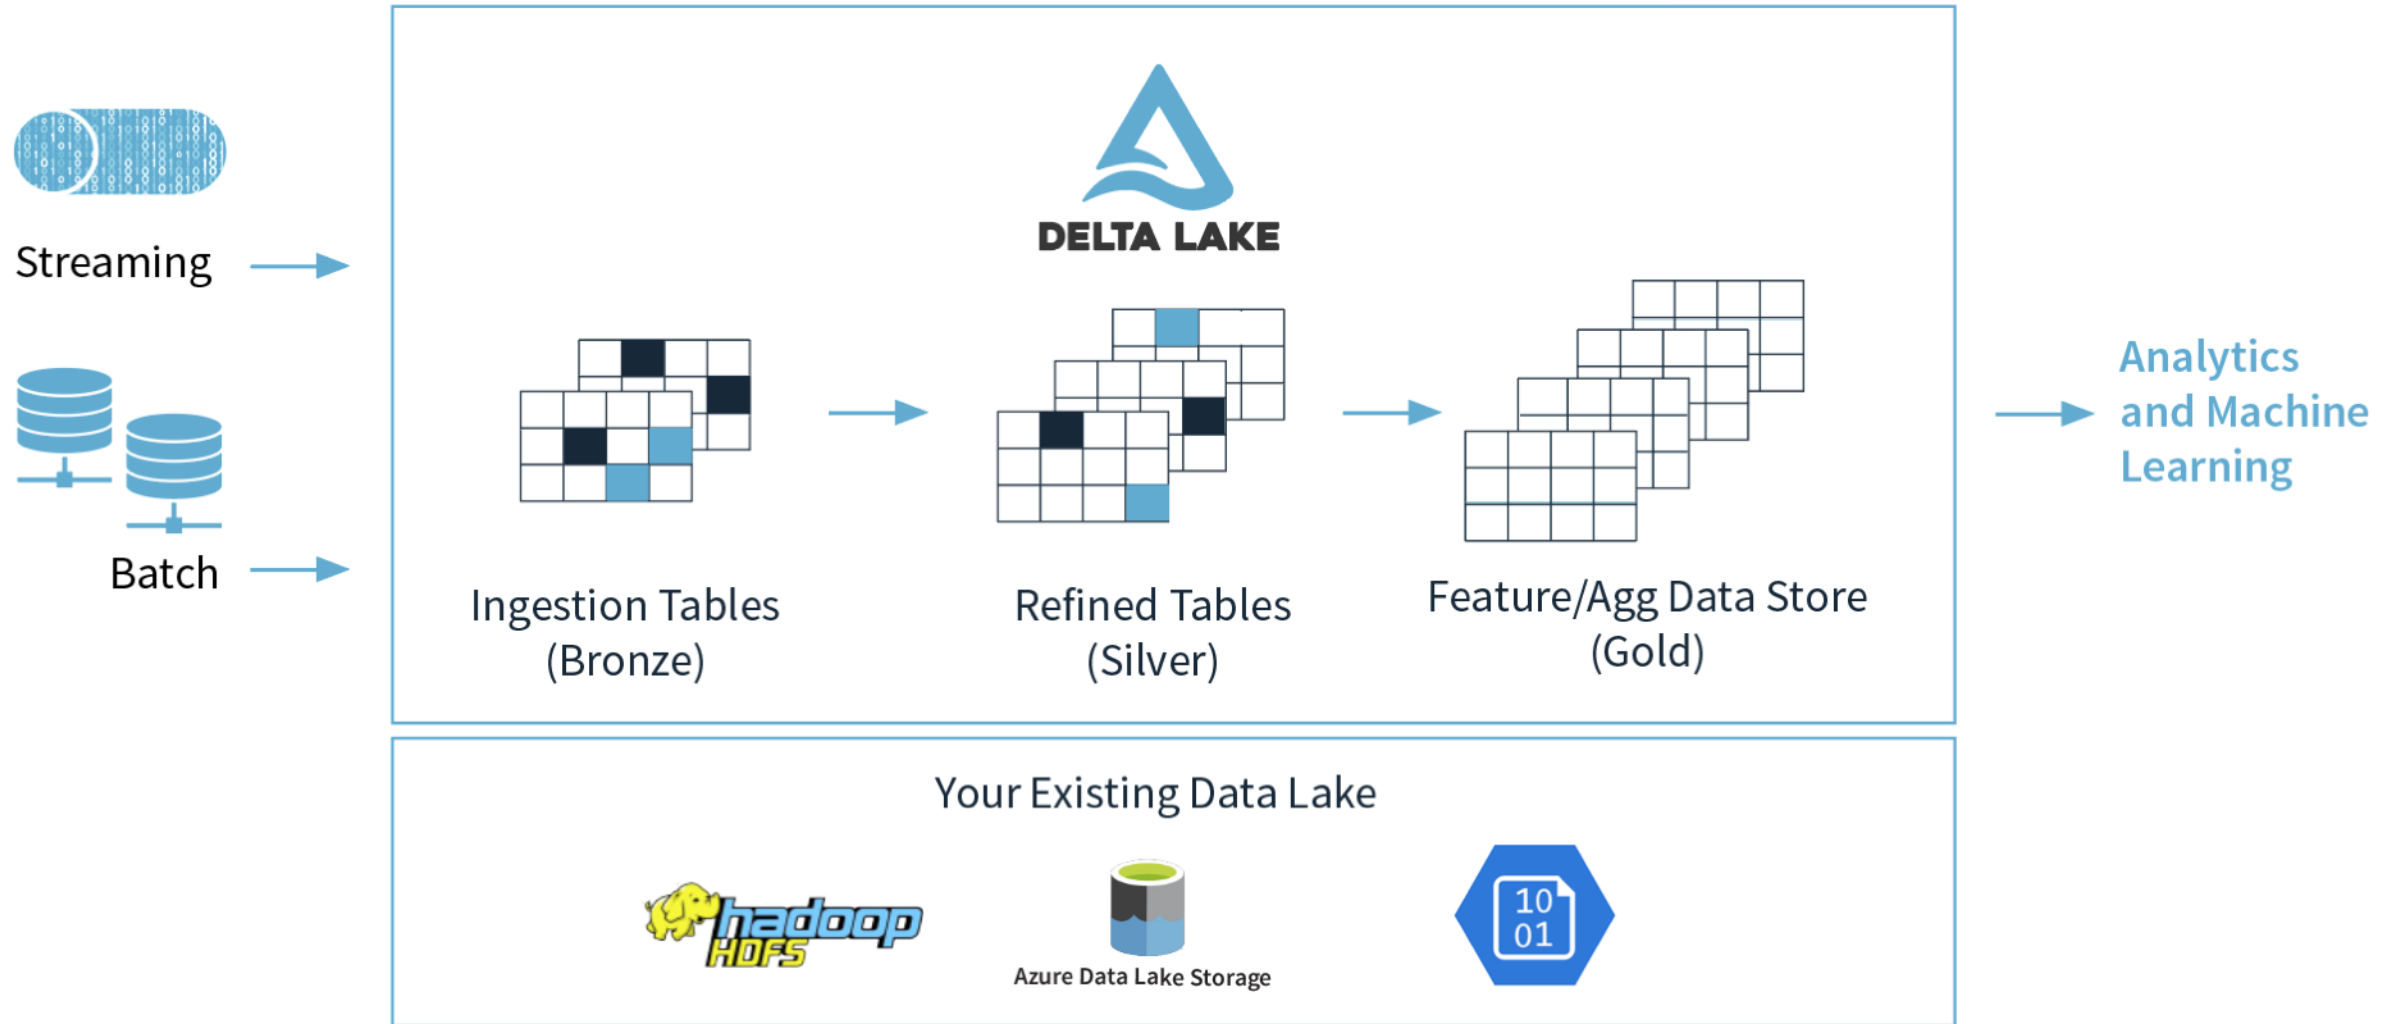
Problems with Parquet format

- No schema enforcement
- Updates – rewrite entire file/all files
- Transactions inconsistency
- No interoperability between batch & streaming workloads
- No versioning

Advantages of Delta Lake over CSV



Medallion Architecture



Delta Lake – Key Features



ACID Transactions




Scalable
Metadata




Time Travel



Open Source




Unified
Batch/Streaming



Schema Evolution
/ Enforcement



Audit History



DML Operations

CONVERT

```
CONVERT TO DELTA database_name.table_name; -- only for Parquet tables
```

```
CONVERT TO DELTA parquet.`s3://my-bucket/path/to/table`
```

```
    PARTITIONED BY (date DATE); -- if the table is partitioned
```

```
-- Uses Iceberg manifest for metadata
```

```
CONVERT TO DELTA iceberg.`s3://my-bucket/path/to/table`;
```

Transaction Log

Transaction Log
Single Commits

(Optional) Partition Directories
Data Files



```
graph LR; A[Transaction Log  
Single Commits] --> B[my_table/]; A --> C[_delta_log/]; A --> D[00000.json]; A --> E[00001.json]; A --> F[date=2019-01-01/]; A --> G[file-1.parquet];
```


my_table/
_delta_log/
00000.json
00001.json
date=2019-01-01/
file-1.parquet

Transaction Log

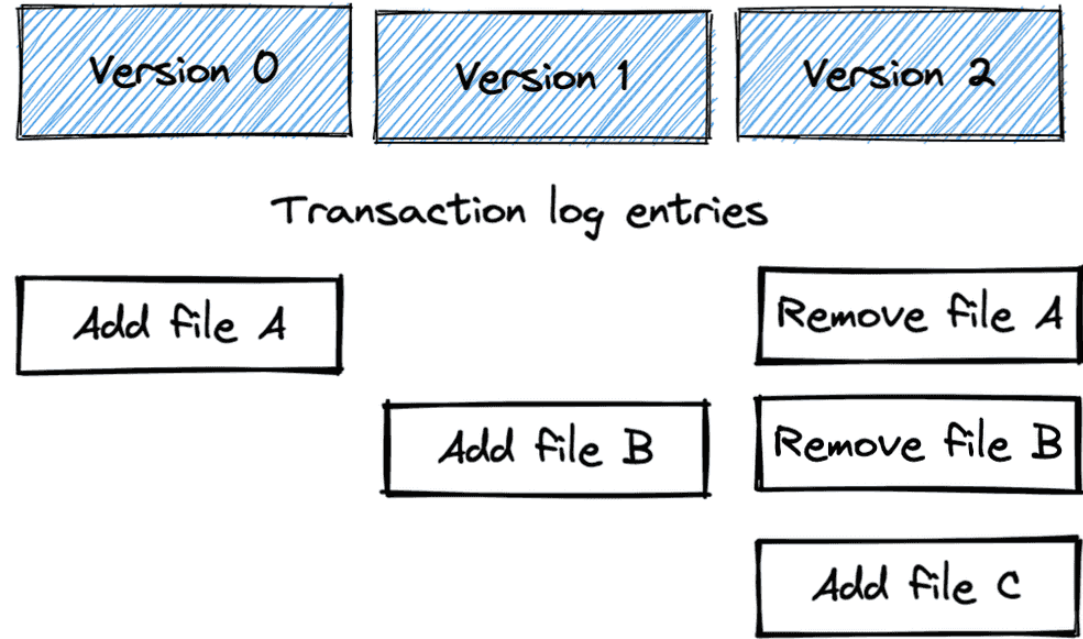


Time Travel

Delta table versions



Delta table transactions



Metadata Tools

Command	Result
DESCRIBE TABLE	Schema of a table
DESCRIBE EXTENDED	Schema of a table + table details
DESCRIBE HISTORY	Table's changes & versions

DESCRIBE TABLE

- DESCRIBE TABLE Lakedb.deltaUsers

	col_name ▲	data_type ▲	comment ▲
1	UserId	int	null
2	DisplayName	string	null
3	LastAccessDate	date	null
4	Reputation	int	null

DESCRIBE TABLE EXTENDED

- DESCRIBE EXTENDED Lakedb.deltaUsers

	col_name ▲	data_type ▲	comment ▲
1	UserId	int	null
2	DisplayName	string	null
3	LastAccessDate	date	null
4	Reputation	int	null
5			
6	# Detailed Table Information		
7	Catalog	spark_catalog	
8	Database	Lakedb	
9	Table	deltaUsers	
10	Type	EXTERNAL	
11	Location	dbfs:/datarelay2023/delta/deltausers	
12	Provider	delta	
13	Owner	root	
14	Table Properties	[delta.minReaderVersion=1,delta.minWriterVersion=2]	

View Table Details

- DESCRIBE DETAIL
‘ /delta/deltausers/ ’
- DESCRIBE DETAIL
Lakedb.deltaUsers

Column	Type	Description
format	string	Format of the table, that is, delta.
id	string	Unique ID of the table.
name	string	Name of the table as defined in the metastore.
description	string	Description of the table.
location	string	Location of the table.
createdAt	timestamp	When the table was created.
lastModified	timestamp	When the table was last modified.
partitionColumns	array of strings	Names of the partition columns if the table is partitioned.
numFiles	long	Number of the files in the latest version of the table.
sizeInBytes	int	The size of the latest snapshot of the table in bytes.

format	id	name	description	location	createdAt	lastModified	partitionColumns
delta	4a67f614-32f2-465f-85fa-5e466983393c	spark_catalog.lakedb.deltausers	null	dbfs:/datarelay2023/delta/deltausers	2023-09-30T01:20:30.202+0000	2023-09-30T01:21:46.000+0000	[]
numFiles	sizeInBytes	properties	minReaderVersion	minWriterVersion	tableFeatures	statistics	
2	2409	{}	1	2	▶ ["appendOnly",	{}	

History

- DESCRIBE HISTORY Lakedb.deltaUsers
- DESCRIBE HISTORY delta.`/delta/users`

	version	timestamp	userId	userName	operation	operationParameters	job
1	2	2023-09-30T01:21:46.000+0000	4809069939003851	kamil@nowinski.net	MERGE	<pre>{ "predicate": "[\"(UserId#14978 = UserId#14982)\"]", "matchedPredicates": "[{\"actionType\": \"update\"}]", "notMatchedPredicates": "[{\"actionType\": \"insert\"}]", "notMatchedBySourcePredicates": [] }</pre>	null
2	1	2023-09-30T01:21:10.000+0000	4809069939003851	kamil@nowinski.net	MERGE	<pre>{ "predicate": "[\"(UserId#13638 = UserId#13646)\"]", "matchedPredicates": "[{\"actionType\": \"update\"}]", "notMatchedPredicates": "[{\"actionType\": \"insert\"}]", "notMatchedBySourcePredicates": [] }</pre>	null

notebook	clusterId	readVersion	isolationLevel	isBlindAppend	operationMetrics	userMetadata
<pre>{ "notebookId": "2353687493752337" }</pre>	0926-221804-4448n59m	1	WriteSerializable	false	<pre>{ "numTargetRowsCopied": "2", "numTargetRowsDeleted": "0", "numTargetFilesAdded": "2", "numTargetBytesAdded": "2409", "numTargetBytesRemoved": "1247", "numTargetDeletionVectorsAdded": "0", "numTargetRowsMatchedUpdated": "1", "executionTimeMs": "8325", "numTargetRowsInserted": "0", "numTargetRowsMatchedDeleted": "0", "scanTimeMs": "1742", "numTargetRowsUpdated": "1", "numOutputRows": "3", "numTargetDeletionVectorsRemoved": "0", "numTargetRowsNotMatchedBySourceUpdated": "0", "numTargetChangeFilesAdded": "0", "numSourceRows": "1", "numTargetFilesRemoved": "1", "numTargetRowsNotMatchedBySourceDeleted": "0", "rewriteTimeMs": "6485" }</pre>	null
<pre>{ "notebookId": "2353687493752337" }</pre>	0926-221804-4448n59m	0	WriteSerializable	false	<pre>{ "numTargetRowsCopied": "0", "numTargetRowsDeleted": "0", "numTargetFilesAdded": "1", "numTargetBytesAdded": "1247", "numTargetBytesRemoved": "0", "numTargetDeletionVectorsAdded": "0", "numTargetRowsMatchedUpdated": "0", "executionTimeMs": "3056", "numTargetRowsInserted": "3", "numTargetRowsMatchedDeleted": "0", "scanTimeMs": "1979", "numTargetRowsUpdated": "0", "numOutputRows": "3", "numTargetDeletionVectorsRemoved": "0", "numTargetRowsNotMatchedBySourceUpdated": "0", "numTargetChangeFilesAdded": "0", "numSourceRows": "3", "numTargetFilesRemoved": "0", "numTargetRowsNotMatchedBySourceDeleted": "0", "rewriteTimeMs": "970" }</pre>	null
<pre>{ "notebookId": "2353687493752337" }</pre>	0926-221804-4448n59m	null	WriteSerializable	true	{}	null

DEMO 1



DEMO Agenda

- Convert CSV to Parquet and to Delta
- Read, insert & update data (DML)
- Transaction Log & DESCRIBE command
- Time Travel

Schema Validation

- All DataFrame **columns** must exist in the target table
- DataFrame column **data types** must match
- DataFrame **column names** cannot differ **only by case**

Schema Evolution

Delta Lake allows for schema evolution

1: Suppose you have this Delta table

first_name	age
bob	47
li	23
leonard	51



2: Data to append

first_name	age	country
frank	68	usa
jordana	26	brasil

→ Extra column of data

3: Schema enforcement will prevent the mismatched append

```
df.write.format("delta").mode("append").save("tmp/fun_people")
```

→ 🚔🛑 Exception: no data appended!

4: Use mergeSchema to enable schema evolution 🥰

```
df.write.option("mergeSchema", "true").mode("append").format("delta").save("tmp/fun_people")
```

```
spark.read.format("delta").load("tmp/fun_people").show()
```

first_name	age	country
jordana	26	brasil
frank	68	usa
leonard	51	null
bob	47	null
li	23	null

→ So amazingly cool!

Schema Evolution - Automatic

Columns that are present in the DataFrame but missing from the table are automatically added as part of a write transaction when:

- write or writeStream have `.option("mergeSchema", "true")`
- `spark.databricks.delta.schema.autoMerge.enabled` is true

Update Table Schema

Columns	Query (in SQL)	Behaviour without schema evolution (default)	Behaviour with schema evolution
Target columns: key, value Source columns: key, value, new_value	<pre>MERGE INTO target_table t USING source_table s ON t.key = s.key WHEN MATCHED THEN UPDATE SET * WHEN NOT MATCHED THEN INSERT *</pre>	The table schema remains unchanged; only columns key, value are updated/inserted.	The table schema is changed to (key, value, new_value). Existing records with matches are updated with the value and new_value in the source. New rows are inserted with the schema (key, value, new_value).
Target columns: key, old_value Source columns: key, new_value	<pre>MERGE INTO target_table t USING source_table s ON t.key = s.key WHEN MATCHED THEN UPDATE SET * WHEN NOT MATCHED THEN INSERT *</pre>	UPDATE and INSERT actions throw an error because the target column old_value is not in the source.	The table schema is changed to (key, old_value, new_value). Existing records with matches are updated with the new_value in the source leaving old_value unchanged. New records are inserted with the specified key, new_value, and NULL for the old_value.
Target columns: key, old_value Source columns: key, new_value	<pre>MERGE INTO target_table t USING source_table s ON t.key = s.key WHEN MATCHED THEN UPDATE SET new_value = s.new_value</pre>	UPDATE throws an error because column new_value does not exist in the target table.	The table schema is changed to (key, old_value, new_value). Existing records with matches are updated with the new_value in the source leaving old_value unchanged, and unmatched records have NULL entered for new_value. See note (1) .
Target columns: key, old_value Source columns: key, new_value	<pre>MERGE INTO target_table t USING source_table s ON t.key = s.key WHEN NOT MATCHED THEN INSERT (key, new_value) VALUES (s.key, s.new_value)</pre>	INSERT throws an error because column new_value does not exist in the target table.	The table schema is changed to (key, old_value, new_value). New records are inserted with the specified key, new_value, and NULL for the old_value. Existing records have NULL entered for new_value leaving old_value unchanged. See note (1) .



DEMO 2: **Schema validation & enforcement**

Release 3.2

- <https://github.com/delta-io/delta/releases/tag/v3.2.0>
- [Support for Apache Spark 3.5.](#)

Release 4.0 (preview)

- <https://github.com/delta-io/delta/releases/>
- Documentation:
<https://github.com/delta-io/delta/releases/tag/v4.0.0rc1>
- Built on top of Apache Spark 4.0.0!

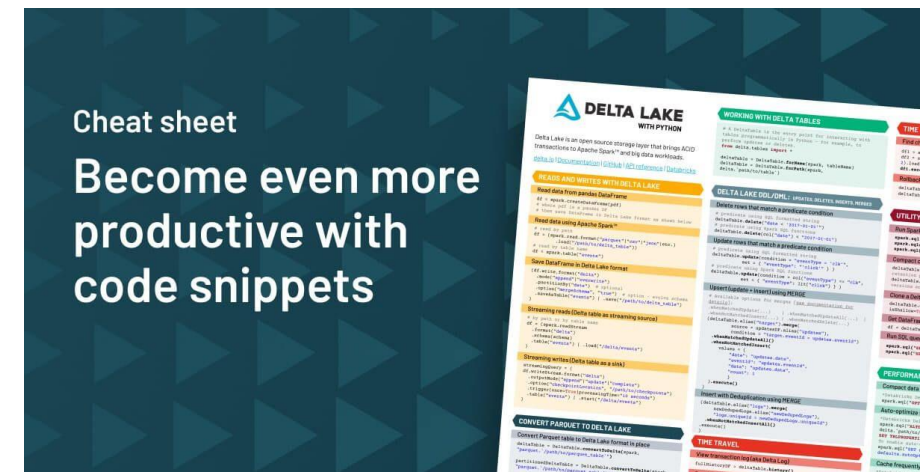
Summary

Summary

- Delta Lake is optimized storage layer
- Delta Lake operates on Parquet files underneath
- File-based transaction log for [ACID transactions](#)
- Open Source
- Supported by MS Fabric, Databricks, Synapse & ADF
- Default storage format for all operations on Azure Databricks

Resources

- <https://delta.io/>
- <https://docs.delta.io/latest/releases.html>
- <https://github.com/delta-io/delta/>
- <https://github.com/delta-io/delta/releases/>
- <https://github.com/delta-io/delta-examples/>
- [Table batch reads and writes](#)
- [Lambda Architecture](#)
- [Delta Lake Blogs](#)
- [What is Delta Lake? \(Microsoft\)](#)
- [Getting Started with Delta Lake \(Databricks\)](#)
- [Follow on LinkedIn: Delta Lake](#)
- [Unpacking The Transaction Log](#)
- [Cheat sheets on Learn with AzurePlayer \(free\)](#)



Resources – cont.

How to engage?



delta.io



delta-users
Slack



Delta Lake
YouTube channel



delta-users
Google Group



Delta Lake
GitHub Issues



Delta Lake RS
Bi-weekly meetings



data-ai-online



Thank you
for your feedback

Thank you!

¡Gracias



kamil@azureplayer.net



@AzurePlayer.bsky.social



<https://AzurePlayer.net/slides>



kamilnowinski



AzurePlayer.net

Kamil Nowinski

Microsoft Data Platform MVP

Analytics Architect, Azure DevOps Engineer Expert