

DISCOVERY DAY 2025

REDEFINING POSSIBILITIES
WITH **SQL SERVER 2025,**
MICROSOFT FABRIC,
AND AI

12.06.2025 | SOFIA

Partners:



Hospitality Partner:



Media Partners:



Food & Drinks Partners:



Gelateria Diana



Host:



MICROSOFT FABRIC: WHERE TO START?

DISCOVERY DAY 2025



KAMIL NOWINSKI

KAMIL NOWINSKI

Group Manager, Analytics Architect
Avanade UK



DISCOVERY
DAY 2025



azureplayer.net



[@KamilNowinski](https://www.youtube.com/@KamilNowinski)



[@AzurePlayer.bsky.social](https://bsky.social/@AzurePlayer)



[linkedin.com/in/KamilNowinski](https://www.linkedin.com/in/KamilNowinski)

FAVORITE STUFF:





Query data in a KQL database in Microsoft Fabric | Lab 12

Kamil Data Geek - Azure explained • 368 views • 2 months ago



Get started with Data Activator in Microsoft Fabric | Lab 11

Kamil Data Geek - Azure explained • 705 views • 2 months ago



Ingest data with Spark and Microsoft Fabric notebooks | Lab 10

Kamil Data Geek - Azure explained • 498 views • 3 months ago



Use real time eventstreams in Microsoft Fabric | Lab 09

Kamil Data Geek - Azure explained • 335 views • 3 months ago

Ask SQL Family ▶ Play all

This set of videos promote episodes of "Ask SQL Family" podcast.



Leila Etaati - You won't be forgotten | Ask SQL Family...

Kamil Data Geek - Azure explained
73 views • 2 months ago



Amit Bansal - Is the role of DBA changing? | Ask SQL...

Kamil Data Geek - Azure explained
37 views • 4 months ago



Guy in a Cube - Become a better presenter | Ask SQL...

Kamil Data Geek - Azure explained
44 views • 5 months ago



Guy in a Cube - how did "YOOO" become a hallmark?...

Kamil Data Geek - Azure explained
55 views • 5 months ago



Pawel Potasinski - What's important in your public...

Kamil Data Geek - Azure explained
84 views • 6 months ago



Simon Whiteley - What is hiding behind the term...

Kamil Data Geek - Azure explained
141 views • 8 months ago

www.AzurePlayer.net/YouTube

Slides available:




<https://azureplayer.net/slides>


Agenda

- Intro to Microsoft Fabric
 - Fabric Workloads
 - Evolution of Microsoft Fabric
- Ingest and store capabilities in Fabric
- Data Factory
- Real-Time Analytics
- Data Warehousing
- CI/CD
- Monitoring


Evolutionary History of Microsoft Fabric




Pre-Fabric




Fabric scope




Fabric artefacts



Fabric experiences



Introduction to Microsoft Fabric





Unified data platform for the era of AI



Data Factory



Data Engineering



Data Warehouse



Data Science



Real-Time Intelligence




Power BI


- **Data Factory:** Data integration combining Power Query with the scale of Azure Data Factory to move and transform data.
- **Data Engineering:** Data engineering with a Spark platform for data transformation at scale.
- **Data Warehouse:** Data warehousing with industry-leading SQL performance and scale to support data use.
- **Data Science:** Data science with Azure Machine Learning and Spark for model training and execution tracking in a scalable environment.
- **Real-Time Intelligence:** Real-Time Intelligence to query and analyze large volumes of data in motion in near real-time.
- **Power BI:** Business intelligence for translating data to decisions.

Automatically store all data in OneLake

- Built on top of Azure Data Lake Storage **(ADLS) Gen2**, OneLake supports any type of file
- All Fabric data items store their data automatically in OneLake in **Delta Parquet** format
- Compatible with existing ADLS Gen2 applications, including Azure **Databricks**
- Data in OneLake is addressable as if it's one big ADLS storage account
- Every workspace appears as a **container** within the storage account
- Azure Databricks **Unity Catalog** integration (in development)




Did I say Delta Parquet?




Bring in data from anywhere

Mirroring



Add or manage data from your cloud data warehouses and databases


Fabric compute engines



Multi-cloud shortcuts


Virtualize data in OneLake from Azure, Dataverse, Google, and Amazon S3 without moving or duplicating it

Understand and manage the Fabric environment



Understand Fabric concepts

- Tenant
- Capacity
- Domain (optional)
- Subdomain (optional)
- Workspace
- Items



What's in a Fabric workspace?

Home

Create

Browse

OneLake

Monitor

Real-Time

Workloads

Workspaces

Workspace 1

...

Data Engineering

Data Engineering

Workspace 1

Development

Create app

Manage access

Workspace settings

+ New item

New folder


Import

Filter by keyword


Filter

Name	Type	Task	Owner	Refreshed	Next refresh	Endorsement	Sensitivity	Included in app
Copilot Demo	Folder	—	—	—	—	—	—	
Direct Lake demo	Folder	—	—	—	—	—	—	
	Folder	—	—	—	—	—	—	
	Folder	—	—	—	—	—	—	
	Folder	—	—	—	—	—	—	
Dataflow 1	Dataflow...	—			N/A	—	—	
env_with_libraries	Environm...	—			—	—	—	
modelRenameTest	Semantic...	—			N/A	—	—	
semantic_link_files	Lakehouse	—			—	—	—	
semantic_link_files	Semantic...	—			N/A	—	—	
semantic_link_fil	SQL anal...	—			N/A	—	—	

Fabric workspace settings

- 
- General
 - License info
 - Azure connections
 - System storage
 - Git integration
 - OneLake
 - Workspace identity
 - Network security
 - Monitoring


- manage Azure connections (e.g., ADLS Gen2, Azure Log Analytics)
- set up Git integration for CI/CD
- configure OneLake settings (including shortcut caching)
- create workspace identities for authentication
- manage network security and private endpoints
- and add Eventhouse for monitoring workspace activity

- 
- Power BI
 - General
 - Data connections
 - Embed codes
 - Delegated Settings
 - OneLake settings
 - Data Engineering/Science
 - Spark settings**
 - Data Factory
 - Apache Airflow runtime settings


- manage Spark configurations
- set a default Spark pool for the workspace
- adjust compute settings for notebooks and Spark jobs
- and configure runtime settings for Apache Airflow jobs

Data security in Fabric

- Workspace roles
- Item permissions
- Compute or granular permissions
- OneLake data access controls



Ingest and store data in Fabric



Flexible and comprehensive ingestion capabilities

-  Dataflow Gen2 (low-code)
-  Data Factory Pipeline
-  Notebook
-  Shortcut
-  Mirroring
-  Eventstream (streaming data)
-  Direct uploads

Diverse storage options



Lakehouse



Data Warehouse




SQL Database



Eventhouse (KQL Database)

Integrate data sources in Fabric



Data integration in Fabric

 Data Factory workload:

 DataflowsGen2

 Pipelines


 Notebooks

 Real-Time Intelligence

Data Pipelines

Pipeline concepts:

- Activities
 - Data transformation
 - Control flow
- Parameters
- Pipeline runs



Fabric notebooks

- Supported languages:
 - SparkSQL, Scala, Python, R, HTML
- Integrated data processing
- Scheduling and automation
- Data exploration and visualization
- ML and AI
- Spark integration

Transform Sales Data - Test | Saved

Search

Home Edit Run View

Comments History Develop Share

Run all Connect PySpark (Python) Environment Workspace default Data Wrangler

Explorer

```
1 # Step 1: Create schema and import data from bronze layer. We'll transform this data using this notebook and load to a delta t
2
3 from pyspark.sql.types import *
4
5 # Create the schema for the table
6 orderSchema = StructType([
7     StructField("SalesOrderNumber", StringType()),
8     StructField("SalesOrderLineNumber", IntegerType()),
9     StructField("OrderDate", DateType()),
10    StructField("CustomerName", StringType()),
11    StructField("Email", StringType()),
12    StructField("Item", StringType()),
13    StructField("Quantity", IntegerType()),
14    StructField("UnitPrice", FloatType()),
15    StructField("Tax", FloatType())
16 ])
17
18 # Import all files from bronze folder of lakehouse
19 df = spark.read.format("csv").option("header", "true").schema(orderSchema).load("Files/Bronze/*.csv")
20
21 # display(df)
```

[1] ✓ - Session ready in 8 sec 610 ms. Command executed in 2 sec 374 ms by Shannon Lindsay on 2:00:07 PM, 8/14/23 PySpark (Python)

```
1 %%sql
2
3 -- Remove this step - using for testing
4
5 -- Drop a table if it exists
6 DROP TABLE IF EXISTS sales_silver
```

[2] ✓ - Command executed in 14 sec 73 ms by Shannon Lindsay on 2:00:22 PM, 8/14/23 Spark SQL

No data available


```
1 # Step 2: Clean up data to load to delta table (silver layer). These steps will fix nulls and add a timestamp for audit purpos
2
3 # Create view to add audit columns (timestamp) - ask: why wouldn't I just use SQL magic here?
4 df.createOrReplaceTempView(f"vw_sales")
5 query = """
6 CREATE TABLE
7     sales_silver USING DELTA
8 AS
9 SELECT
10 *,
11 'FILE_NAME' AS FILE_NAME,
```

Not connected AutoSave: On

Selected Cell 1 of 8 cells

Play with data & have fun!

Integrate data using Data Factory






Data pipeline

Ingest and orchestrate activities at scale

Familiar authoring canvas experience



OneLake



Copy job

Easily ingest and move data at scale

Supports both batch and near real-time incremental copy (CDC)

Data factory Search (Alt + Q) Trial: 18 days left

New items saved to: Contoso workspace Show less ^

New

- Dataflow Gen2 (Preview)
- Data pipeline
- Copy job
- Apache Airflow project
- Data Factory mount

Recommended

- Not sure where to start? Tell Copilot what you want to do and it will do it for you. Launch Dataflow Copilot
- Not sure where to start? Tell Copilot what you want to do and it will do it for you. Launch Pipelines Copilot
- Introduction to data integration Essential concepts to know. Watch video
- Getting started with dataflow Learn the basics of dataflow. Watch the tutorial video
- Getting started with data pipelines Learn the basics of data pipelines. Watch the tutorial video

Quick access

Recent Favorites My apps Filter by keyword Filter v

Name	Type	Opened	Owner	Endorsement	Sensitivity	Workspace
Cool data mart	Datamart	7m ago	Tim Deboar	—	General	Contoso workspace
Data flow for triggers	Data flow	7m ago	Tim Deboar	—	—	Contoso workspace
User data	Datamart	7m ago	Tim Deboar	Certified	—	Contoso workspace
Copy data pipeline	Pipeline	7m ago	Tim Deboar	—	—	Contoso workspace



OneLake



Data pipeline

Ingest and orchestrate activities at scale

Empower every person to integrate data

The screenshot displays the Azure Data Factory user interface. At the top, there's a navigation bar with tabs for 'Home', 'Activities', 'Run', and 'View'. Below this is a toolbar with icons for 'Validate', 'Run', 'Schedule', 'View run history', 'Copy data', 'Dataflow', 'Notebook', 'Lookup', and 'Invoke pipeline'. The main workspace shows a pipeline with two activities: 'Copy Products' (under the 'Copy data' category) and 'Send Completion Email' (under the 'Office 365 Outlook (Preview)' category). Both activities have a green checkmark indicating they are configured correctly. The bottom of the interface has tabs for 'Parameters', 'Variables', 'Settings', and 'Output', with a '+ New' button under the 'Variables' tab. A red mouse cursor is pointing at the bottom center of the workspace.



OneLake




Run and monitor pipelines

1. Validate
2. Run
3. Schedule
4. View run history

The screenshot displays the Azure Data Studio interface. The top navigation bar includes 'Home', 'Activities', 'Run', and 'View'. The 'Run' tab is active, showing a toolbar with 'Validate', 'Run', 'Schedule', and 'View run history' buttons. A red box highlights the 'Validate', 'Run', and 'Schedule' buttons. The main workspace shows a 'Copy data' task configuration with the name 'Copy from To process to In process'. The bottom panel shows the 'General' tab of the task configuration, with fields for 'Name', 'Description', 'Timeout', and 'Retry'. The right sidebar displays a table of pipeline runs.


Name	Status	Start time
Process rides pipeline	Completed	3:17 PM, 6/21/23
Process rides pipeline	Failed	2:31 PM, 6/28/23
Process rides pipeline	Completed	2:35 PM, 6/28/23
Process rides pipeline	Completed	2:44 PM, 6/28/23
Process rides pipeline	Completed	2:50 PM, 6/28/23
Process rides pipeline	Completed	3:41 PM, 6/28/23
Process rides pipeline	Completed	3:50 PM, 6/28/23
Process rides pipeline	Completed	3:33 PM, 6/30/23
Process rides pipeline	Completed	3:55 PM, 6/30/23

Integrate data using Notebooks



Notebooks in Fabric

- Use notebooks to:
 - Ingest and transform data
 - Schedule and orchestrate (used with a pipeline)
 - Manage dependencies
 - Implement error handling and monitoring
 - Automate workflows
- Trigger a notebook run by:
 - Interactive run
 - Run as pipeline activity
 - Scheduler



The screenshot shows a Jupyter Notebook interface with a toolbar at the top containing icons for navigation and actions like 'Add code cell', 'Clear all output(s)', and 'Find and replace'. The notebook title is 'Getting started with your Notebook'. It contains three code cells with the following Python code:

```
1 # Use the 2 magic commands below to reload the modules if your module has updates during the cu
2 # %load_ext autoreload

1 # %autoreload 2
2
3 import builtin.Code.SampleModule as SampleModule
4 # Now use the exported members from this module with identifier: `SampleModule`
5 dir(SampleModule)
6

1 df = spark.read.parquet("Files/SampleData.parquet")
2
3 display(df)
```

Run notebooks in pipelines

HomeActivitiesRunView



 Validate  Run  Schedule  Add trigger (preview)  View run history  Copy data  Dataflow  Notebook 

Copy data

 Copy from Bronze to Silver

Notebook

 Run notebook to transform files

GeneralSettings

 Please review this item carefully before adding it to the pipeline, as others in your organization may have access to notebooks in this workspace.

Workspace *

Learn



 Refresh

Notebook *

Transform Sales Data - Test




 Refresh  Open  New

> Base parameters

> Advanced settings


Real-Time Analytics in Fabric

<https://aka.ms/mslearn-get-started-real-time-fabric>



Fabric Real-Time Analytics

- **Eventstream:** Real-time data stream ingestion and transformation
- **Eventhouse:** Scalable storage and querying of real-time data
- **KQL Queryset:** A tool for creating and running queries in the KQL databases within an eventhouse
- **Real-Time Dashboard:** Interactive visualization and exploration of real-time data
- **Activator:** Automated actions triggered by real-time events



Ingest and transform real-time data

The screenshot displays the Azure Real-Time Intelligence (RTI) interface in a web browser. The top navigation bar shows the workspace name 'stock-stream' and a search bar. The main canvas is in 'Edit mode', showing a data pipeline with three components: a 'stock' source, a 'stock-stream' connector (containing 'stock-stream-stream'), and a 'Transform events or add destination' sink. Below the canvas, the 'Test result' tab is active, showing a table of test data for the 'stock-stream' workspace. The table includes columns for time, symbol, sector, securityType, bidPrice, and bidSize. The data is filtered for the time range '12/03/24 09:52:03 AM - 12/03/24 10:52:03 AM'.


Test result Authoring errors

Data format: Json Time range: 12/03/24 09:52:03 AM - 12/03/24 10:52:03 AM

time	symbol	sector	securityType	bidPrice	bidSize
2024-12-03T18:52:04.241Z	BMZM	retailing	commonstock	2316.84	41
2024-12-03T18:52:04.241Z	NSFT	softwareservices	commonstock	350.63	3
2024-12-03T18:52:04.241Z	HOOJ	mediaentertainment	commonstock	1250.28	31
2024-12-03T18:52:04.241Z	BMZM	retailing	commonstock	2306.84	123
2024-12-03T18:52:04.241Z	NSFT	softwareservices	commonstock	370.63	23
2024-12-03T18:52:04.241Z	HOOJ	mediaentertainment	commonstock	1370.28	66
2024-12-03T18:52:04.241Z	BMZM	retailing	commonstock	2316.84	41
2024-12-03T18:52:04.241Z	NSFT	softwareservices	commonstock	330.63	89
2024-12-03T18:52:04.241Z	HOOJ	mediaentertainment	commonstock	1330.28	103


Store and query real-time data

- Scalable Storage
- KQL Databases
- KQL Query Language
- KQL Querysets
- AI Assistance




Visualizing Real-Time Data

Real-time Dashboard




Power BI Dashboard



Automate actions with Activator

- Each record in a stream of data represents an **event**
- Model real-world or conceptual **objects** based on data values
- Define **rules** based on conditions to trigger actions
- Use field data values to define object **properties** and reference them from multiple rules



Use KQL to query tables

Filter columns

- Use the **project** keyword
- Retrieve only needed columns
- **Example:** project trip_id, vendor_id

Filter rows

- Use the **where** keyword
- Filter rows based on criteria
- **Example:**
where getmonth(pickup_datetime)

Summarize data

- Summarize data for patterns
- Use functions like avg and by
- **Example:**
summarize avg(fare_amount)

Advanced KQL features


Materialized views

- Summary of data
- Real-time updates
- Query like tables
- Efficient analysis

Stored functions

- Encapsulate queries
- Parameterized
- Reusable
- Efficient

Understand data warehouse fundamentals




Understand data warehouse fundamentals

The process of building a modern data warehouse typically consists of:

- **Data ingestion** - moving data from source systems into a data warehouse.
- **Data storage** - storing the data in a format that is optimized for analytics.
- **Data processing** - transforming the data into a format that is ready for consumption by analytical tools.
- **Data analysis and delivery** - analyzing the data to gain insights and delivering those insights to the business.

Data warehouse basics


Ingest



Store



Process




Analyze



Data warehouse design

Tables organized into:

- **Fact tables** for numeric data
- **Dimension tables** for descriptive info





Schema design

★ Star

vs.




❄️ Snowflake



FROM RELATIONAL TO STAR SCHEMA

RELATIONAL SCHEMAS ARE OPTIMIZED FOR NORMALIZATION (=REDUCE DATA REDUNDANCY) AND SPEED OF EXECUTION FOR TRANSACTIONAL QUERIES

STAR SCHEMAS ARE OPTIMIZED FOR ANALYTICAL QUERY PERFORMANCES



PK: PRIMARY KEY. FK: FOREIGN KEY.

NOTE: THIS SCHEMA IS SIMPLISTIC ON PURPOSE.
(E.G. A BOOK COULD BE AUTHORED BY MULTIPLE AUTHORS)

HOW TO MOVE FROM RELATIONAL TO STAR SCHEMA?

1. IDENTIFY THE BUSINESS PROCESS
2. DETERMINE THE GRAIN
3. IDENTIFY THE DIMENSIONS
4. DESIGN THE FACT TABLE
5. DESIGN DIMENSION TABLES
6. HANDLE HIERARCHIES
7. GENERATE SURROGATE KEYS
8. MANAGE SLOWLY CHANGING DIMENSIONS
9. OPTIMIZE, TEST, DOCUMENT

FactBook

FactBookId(SK)
bookId (AK)
title
DimDateId (SK)
publishedDate
DimAuthorId (SK)
authorName
authorFirstName
authorLastName
DimPublisherId (SK)
publisherName
publisherCountry

FACT TABLES CONTAINS THE MEASUREMENT OF "FACT" THAT YOU WANT TO ANALYZE

SURROGATE KEYS (SK) IS A UNIQUE IDENTIFIER FOR EACH RECORD IN THE DATAWAREHOUSE. ACTS AS PRIMARY KEY

ALTERNATE KEYS (AK) MAINTAIN A CONNECTION TO THE ORIGINAL DATA SOURCES

MOST OF STAR SCHEMA HAVE A "DATE" DIMENSION THAT IS COMMON ACROSS ALL ENTITIES

DimDate

DimDateKey
FullDate
DayOfWeek
DayOfMonth
MonthNumber
MonthName
Quarter
Year

DimPublisher

DimPublisherId (SK)
publisherId (AK)
publisherName
publisherCountry

DIMENSION TABLES CONTAINS ATTRIBUTES THAT ARE USED TO FILTER, GROUP OR LABEL DATA CONTAINED IN FACT TABLES

DimAuthor


DimAuthorID (SK)
authorId (AK)
authorName
firstName
lastName

Tables in a data warehouse

Tables in a data warehouse are typically organized in a way that supports efficient and effective analysis of large amounts of data.


- A ***surrogate key*** is a unique identifier for each row in the dimension table. It's often an integer value that is automatically generated by the database management system when a new row is inserted into the table.
- An ***alternate key*** is often a natural or business key that identifies a specific instance of an entity in the transactional source system - such as a product code or a customer ID.

Data warehousing in Fabric



Data warehouse basics

- Massively Parallel (MPP)
- "Serverless"
- Stored in OneLake
- Parquet file format
- T-SQL



UNLIMITED	DATA VOLUME	UNLIMITED
UNSTRUCTURED, SEMI-STRUCTURED, STRUCTURED	TYPE OF DATA	STRUCTURED (TABLES, COLUMNS)
DATA ENGINEER, DATA SCIENTIST	WHO IS USING IT?	DATA WAREHOUSE DEVELOPER, SQL ENGINEER
SPARK (SCALA, PYSPARK, SPARK SQL, R) SPARK NOTEBOOKS, SPARK JOB DEFINITIONS	PRIMARY SKILLSET	SQL SQL SCRIPTS

LAKEHOUSE VS WAREHOUSE


SPARK & T-SQL (SQL ENDPOINT FOR LAKEHOUSE, LIMITED T-SQL COVERAGE)	HOW TO READ DATA?	SPARK & T-SQL (GREAT T-SQL COVERAGE)
✗ No	MULTI-TABLE TRANSACTIONS?	YES ✓
YES, FILES AND TABLES CAN BE SOURCES	SHORTCUTS SUPPORT	YES – ACCESS DATA THROUGH LAKEHOUSE, TABLES CAN BE SOURCES
ACROSS LAKEHOUSE AND WAREHOUSE TABLES & ACROSS LAKEHOUSES	QUERY ACROSS ITEMS	ACROSS LAKEHOUSE AND WAREHOUSE TABLES

Lakehouse + warehouse = better together


Medallion Architecture



Implement continuous integration and continuous delivery (CI/CD)



CI/CD in Fabric



GA (31/03/2025)

Terraform

FabricTools



FabricTools is a PowerShell module to able to do more with Microsoft Fabric and Power BI. It allows for various administrative tasks to be automated and integrated into workflows. Currently at an early stage of development and the module is in its **Public PREVIEW**.

<https://github.com/dataplat/FabricTools>

Implement version control with Git integration

Contained in workspaces

- Azure DevOps or GitHub
- Authenticate
- Connect repository and branch
- Connect and sync

Workspace settings

General

License info

Azure connections

System storage

Git integration

OneLake

Workspace identity

Network security

Power BI

Extension API playground

Data Engineering/Science

Data Factory

Git integration (Preview)

Connect to Git to manage your code and back up your work. [Learn more](#)

Connect Git provider and account

Provider

GitHub

Lz - repo 1

Log out

Connect Git repository and branch

Repository URL

<https://github.com/liavZayde/TestGitIntegration>

Branch * ⓘ

Branch

Git folder ⓘ

Enter name of folder

Connect and sync

Cancel



Azure Player
Play with data & have fun!

CI/CD with deployment pipelines

The screenshot displays the Power BI Deployment Pipelines interface. The top navigation bar includes the Power BI logo, the text "Deployment pipelines", a search bar, and icons for notifications, settings, download, help, and user profile. The left sidebar contains navigation options: Home, Create, Deployment pipelines (selected), Workspaces, My workspace, Demo Pipeline, and a Power BI logo at the bottom.

The main content area shows a "Demo Pipeline" with three stages: Development, Test, and Production. Each stage has a header bar with a rocket icon and a title. Below each header is a card representing the stage's configuration. The cards are connected by green circular arrows with checkmarks, indicating a successful flow.

Development Stage:

- Region: West US 3
- Warehouses: 0
- Copy jobs: 0
- Buttons: Show more, Deploy

Test Stage:

- Region: West US 3
- Warehouses: 0
- Copy jobs: 0
- Buttons: Show more, Deploy

Production Stage:

- Region: West US 3
- Warehouses: 0
- Copy jobs: 0
- Buttons: Show more, Deploy

At the top right of the main content area, there are links for "Deployment history", "Manage Access", "Settings", and a menu icon.

Use deployment pipelines *with* Git

Power BI

Deployment pipelines

Search

Home

Deployment pipelines

Workspaces

My workspace

Pipeline

Demo Pipeline

Power BI

Demo Pipeline

Deployment history

Manage Access

Settings

Development

Test

Production

Development

Region: West US 3

0

1

0

Org appsData pipeli...Notebooks

Show more

Deploy

Test

Deployed: 07/11/2024

Region: West US 3

0

1

0

Org appsData pipeli...Notebooks

Show more

Deploy

Production

Deployed: 07/11/2024

Region: West US 3

0

1

0

Org appsData pipeli...Notebooks

Show more

Deploy

Azure Player

Play with data & have fun!

Automate CI/CD with Fabric REST APIs

Commit

- Commit changes

Update

- Update the workspace

Review

- Review pending and uncommitted changes


List

- List items in the deployment pipeline stage

Deploy

- Deploy stage content

fabric-cicd



Python library designed for use with Microsoft Fabric workspaces. This library supports code-first Continuous Integration / Continuous Deployment (CI/CD) automations to seamlessly integrate Source Controlled workspaces into a deployment framework.


The goal is to assist CI/CD developers who prefer not to interact directly with the Microsoft Fabric APIs.

<https://microsoft.github.io/fabric-cicd/latest/>


Supported items:

- DataPipeline
- Environment
- Notebook
- Report
- SemanticModel
- Lakehouse
- MirroredDatabase
- VariableLibrary
- CopyJob
- Eventhouse
- KQLDatabase
- KQLQueryset
- Reflex
- Eventstream
- Warehouse
- SQLDatabase

fabric-cicd



Monitor analytics solutions



Fabric monitoring hub

View details

- Activity status
- Start and end time
- Duration
- Activity statistics

Monitor

View and track the status of the activities across all the workspaces for which you have permissions within Microsoft Fabric.

Refresh

Filter by keyword

Filter

Column Options

Clear all


To apply filters, select the values from the Filter dropdown menu.

Activity name	Status	Item type	Start time	Submitted by	Location
Purview Hub	Succeeded	Semantic model	12/12/2024, 2:44 PM	Admin Monitoring	Adm
Feature Usage and Adoption	Succeeded	Semantic model	12/12/2024, 2:44 PM	Admin Monitoring	Adm
LHDP0000	Succeeded	Semantic model	12/11/2024, 11:59 AM	System Administrator	fab_v
semantic model refresh	Succeeded	Data pipeline	12/11/2024, 11:58 AM	System Administrator	fab_v
GetProductData	Succeeded	Dataflow Gen2	12/09/2024, 4:08 PM	System Administrator	fab_v
Query products_305df121-cd8f-49a2-8c7a-2203...	Succeeded	Notebook	12/09/2024, 4:03 PM	System Administrator	fab_v
pipeline3	Succeeded	Data pipeline	12/05/2024, 5:20 PM	System Administrator	Dem

Additional monitoring capabilities for admins

Monitor your environment

Monitor capacity usage





Fabric Unified Admin Monitoring (FUAM)

It provides all the items what you need to boost your platform monitoring

"Seamless tenant monitoring experience for Fabric Admins."



"Guided analytical pathways to find the root-cause within minutes."



"Unified data store for your platform logs."



"Enrich FUAM Lakehouse with your data via Shortcuts."

FUAM Lakehouse

Data collection via Fabric Pipelines & Notebooks



Active Items



Activity Logs



Capacity Properties



Capacity Refreshables



Capacity Metrics



Inventory data




Optimization Module




Tenant settings


FUAM | Fabric Unified Admin Monitoring



FUAM | Fabric Unified Admin Monitoring



FUAM | Fabric Unified Admin Monitoring



THANK YOU!
We'd love to hear your thoughts!

Please take a minute to complete short feedback survey.



QR CODE